

Serial HDF5

Matthieu Haefele, CNRS/UPPA

Outline

- HDF5 in the context of Input/Output (IO)
- HDF5 Application Programming Interface (API)
- Playing with Dataspace
- Hands on session

IO in a nutshell

Doing Input / Output is about **TRANSPORTING**

Data stored in **memory**

to / from

Data stored on **disk**




IO in a nustshell

Three criteria / metrics to balance

- Code development / maintenance time
- Performance
- Post-processing requirement

Hardware/Software stack



High level I/O libraries

- Meshes of various complexity (rectilinear, curvilinear, unstructured...)
- Discretized functions on such meshes
- Materials
- ...

Until now, these libraries are mainly used in the context of visualization

I/O libraries

Purpose of I/O libraries:

- Efficient I/O
- Portable binary files
- Higher level of abstraction for the developer

Two main existing libraries:

- Hierarchical Data Format: HDF5
- Network Common Data Form: NetCDF
- Zarr files ?

HDF5 library

HDF5 file:

- HDF5 group: a grouping structure containing instances of zero or more groups or datasets
- HDF5 dataset: a multidimensional array of data elements

HDF5 dataset \Leftrightarrow multidimensional array:

- Name
- Datatype (Atomic, Composite)
- Dataspace (rank, sizes, max sizes) **SIMPLE!**
- Storage layout (contiguous, compact, chunked)

HDF5 High Level APIs

- **Dimension Scale** (H5DS): Enables to attach dataset dimension to scales
- **Lite** (H5LT): Enables to write simple dataset in one call
- **Image** (H5IM): Enables to write images in one call
- **Table** (H5TB): Hides the compound types needed for writing tables
- **Packet Table** (H5PT): Almost H5TB but without record insertion/deletion but supports variable length records
- ...

HDF5 low level API

- **H5F**: File manipulation routines
- **H5G**: Group manipulation routines
- **H5S**: Dataspace manipulation routines
- **H5D**: Dataset manipulation routines
- ...

Just have a look at the outstanding on-line reference manual for HDF5 !

C order versus Fortran order

/* C language */

#define NX 4

#define NY 3

int x,y;

int f[NY][NX];

for (y=0;y<NY;y++)

for (x=0;x<NX;x++)

f[y][x] = x+y;

! Fortran language

integer, parameter :: NX=4

integer, parameter :: NY=3

integer :: x,y

integer, dimension(NX,NY) :: f

do y=1,NY

do x=1,NX

f(x,y) = (x-1) + (y-1)

enddo

enddo

0 1 2 3 1 2 3 4 2 3 4 5

HDF5 first example

```
#define NX      5
#define NY      6
#define RANK    2

int main (void)
{
    hid_t      file, dataset, dataspace;
    hsize_t    dims[2];
    herr_t     status;
    int        data[NY][NX];

    init(data);
    file = H5Fcreate("example.h5", H5F_ACC_TRUNC, H5P_DEFAULT, \
                    H5P_DEFAULT);

    dims[0] = NY;
    dims[1] = NX;
```

HDF5 first example cont.

```
dataspace = H5Screate_simple(RANK, dimsf, NULL);

dataset = H5Dcreate(file, "IntArray", H5T_NATIVE_INT, \
    dataspace, H5P_DEFAULT, H5P_DEFAULT, H5P_DEFAULT);

status = H5Dwrite(dataset, H5T_NATIVE_INT, H5S_ALL, \
    H5S_ALL, H5P_DEFAULT, data);

H5Sclose(dataspace);
H5Dclose(dataset);
H5Fclose(file);

return 0;
}
```

HDF5 high level example cont.

```
status = H5LTmake_dataset_int(file, "IntArray", RANK, \
                                dims, data);

H5Fclose(file);

return 0;
}
```

Variable C type

```
hid_t    file, dataset, dataspace;  
hsize_t  dims[2];  
herr_t   status;
```

- hid_t: handler for any HDF5 objects (file, groups, dataset, dataspace, datatypes...)
- hsize_t: C type used for number of elements of a dataset (in each dimension)
- herr_t: C type used for getting error status of HDF5 functions

File creation

```
file = H5Fcreate("example.h5", H5F_ACC_TRUNC, H5P_DEFAULT, \
                H5P_DEFAULT);
```

- "example.h5": file name
- H5F_ACC_TRUNC: File creation and suppress it if it exists already
- H5P_DEFAULT: file creation property list
- H5P_DEFAULT: file access property list (needed for MPI-IO)

Dataspace creation

```
dimsf[0] = NY;  
dimsf[1] = NX;  
dataspace = H5Screate_simple(RANK, dimsf, NULL);
```

- RANK: dataset dimensionality
- dimsf: size of the dataspace in each dimension
- NULL: specify max size of the dataset being fixed to the size

Dataset creation

```
dataset = H5Dcreate(file, "IntArray", H5T_NATIVE_INT, \
                    dataspace, H5P_DEFAULT, H5P_DEFAULT, H5P_DEFAULT);
```

- file: HDF5 objects where to create the dataset. Should be a file or a group.
- "IntArray": dataset name
- H5T_NATIVE_INT: type of the data the dataset will contain
- dataspace: size of the dataset
- H5P_DEFAULT: default option for property list.

Datatype

- Predefined Datatypes: created by HDF5.
- Derived Datatypes: created or derived from the predefined data types.

There are two types of predefined datatypes:

- **STANDARD:** H5T_IEEE_F32BE means IEEE representation of 32 bit floating point number in big endian.
- **NATIVE:** Alias to standard data types according to the platform where the program is compiled. Ex: on an Intel based PC, H5T_NATIVE_INT is aliased to the standard predefined type, H5T_STD_32LE.

Datatype cont.

A data type can be:

- ATOMIC: cannot be decomposed into smaller data type units at the API level. Ex: integer
- COMPOSITE: An aggregation of one or more data types. Ex: compound data type, array, enumeration

Dataset writing

```
status = H5Dwrite(dataset, H5T_NATIVE_INT, H5S_ALL, \
                  H5S_ALL, H5P_DEFAULT, data);
```

- dataset: HDF5 objects representing the dataset to write
- H5T_NATIVE_INT: Type of the data in memory
- H5S_ALL: dataspace specifying the portion of memory that needs be read (in order to be written)
- H5S_ALL: dataspace specifying the portion of the file dataset that needs to be written
- H5P_DEFAULT: default option for property list (needed for MPI-IO).
- data: buffer containing the data to write

Closing HDF5 objects

```
H5Sclose(dataspace);  
H5Dclose(dataset);  
H5Fclose(file);
```

Opened/created HDF5 objects are closed.

Some comments

```
status = H5Lmake_dataset_int(file, "IntArray", RANK, \
                             dims, data);

H5Fclose(file);

return 0;
}
```

This example is as simple as a **fwrite**, but:

- The generated file is portable
- The generated file can be accessed with HDF5 tools
- Attributes can be added on datasets or groups
- The type of the data can be fixed
- The storage layout can be modified
- Portion of the dataset can be written
- ...

Concept of start, stride, count block

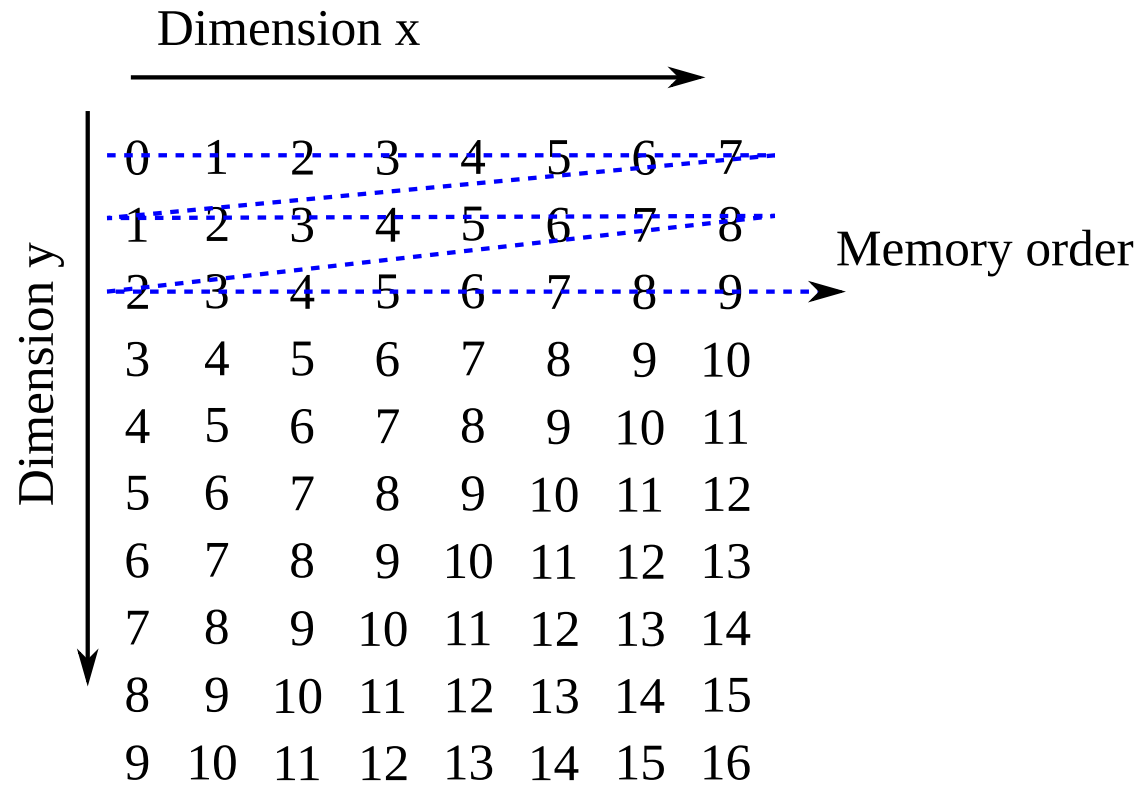
- **start**: Starting location for the hyperslab (default 0)
- **stride**: The number of elements to separate each element or block to be selected (default 1)
- **count**: The number of elements or blocks to select along each dimension
- **block**: The size of the block (default 1)

Conventions for the examples

We consider:

- A 2D array $f[N_y][N_x]$ with $N_x = 8, N_y = 10$
- Dimension x is the dimension contiguous in memory
- Graphically, the x dimension is represented horizontal
- Language C convention is used for indexing the dimensions
- Dimension y is index=0
- Dimension x is index=1


Graphical representation



```
int start[2], stride[2], count[2], block[2];  
start[0] = 0; start[1] = 0;  
stride[0] = 1; stride[1] = 1;  
block[0] = 1; block[1] = 1;
```

Illustration for count parameter


	Dimension x →							
Dimension y ↓	0	1	2	3	4	5	6	7
	1	2	3	4	5	6	7	8
	2	3	4	5	6	7	8	9
	3	4	5	6	7	8	9	10
	4	5	6	7	8	9	10	11
	5	6	7	8	9	10	11	12
	6	7	8	9	10	11	12	13
	7	8	9	10	11	12	13	14
	8	9	10	11	12	13	14	15
	9	10	11	12	13	14	15	16



```
count[0] = 3;  count[1] = 4;
```

Illustration for start parameter


	Dimension x →							
Dimension y ↓	0	1	2	3	4	5	6	7
	1	2	3	4	5	6	7	8
	2	3	4	5	6	7	8	9
	3	4	5	6	7	8	9	10
	4	5	6	7	8	9	10	11
	5	6	7	8	9	10	11	12
	6	7	8	9	10	11	12	13
	7	8	9	10	11	12	13	14
	8	9	10	11	12	13	14	15
	9	10	11	12	13	14	15	16



```
start[0] = 1;  start[1] = 2;  
count[0] = 3;  count[1] = 4;
```

Illustration for stride parameter

	Dimension x →							
Dimension y ↓	0	1	2	3	4	5	6	7
	1	2	3	4	5	6	7	8
	2	3	4	5	6	7	8	9
	3	4	5	6	7	8	9	10
	4	5	6	7	8	9	10	11
	5	6	7	8	9	10	11	12
	6	7	8	9	10	11	12	13
	7	8	9	10	11	12	13	14
	8	9	10	11	12	13	14	15
	9	10	11	12	13	14	15	16



```
start[0] = 1;  start[1] = 2;  
count[0] = 3;  count[1] = 4;  
stride[0] = 3; stride[1] = 1;
```

Illustration for stride parameter

Dimension x →

Dimension y ↓

0	1	2	3	4	5	6	7
1	2	3	4	5	6	7	8
2	3	4	5	6	7	8	9
3	4	5	6	7	8	9	10
4	5	6	7	8	9	10	11
5	6	7	8	9	10	11	12
6	7	8	9	10	11	12	13
7	8	9	10	11	12	13	14
8	9	10	11	12	13	14	15
9	10	11	12	13	14	15	16

y=0 y=1 y=2
↔ ↔ ↔
3 6 | 6 9 | 9 12


```
start[0] = 1;  start[1] = 2;  
count[0] = 3;  count[1] = 2;  
stride[0] = 3; stride[1] = 3;
```

Illustration for block parameter

Dimension x →

Dimension y ↓

0	1	2	3	4	5	6	7
1	2	3	4	5	6	7	8
2	3	4	5	6	7	8	9
3	4	5	6	7	8	9	10
4	5	6	7	8	9	10	11
5	6	7	8	9	10	11	12
6	7	8	9	10	11	12	13
7	8	9	10	11	12	13	14
8	9	10	11	12	13	14	15
9	10	11	12	13	14	15	16



```
start[0] = 1;  start[1] = 2;  
count[0] = 3;  count[1] = 2;  
stride[0] = 3; stride[1] = 3;  
block[0] = 2;  block[1] = 2;
```

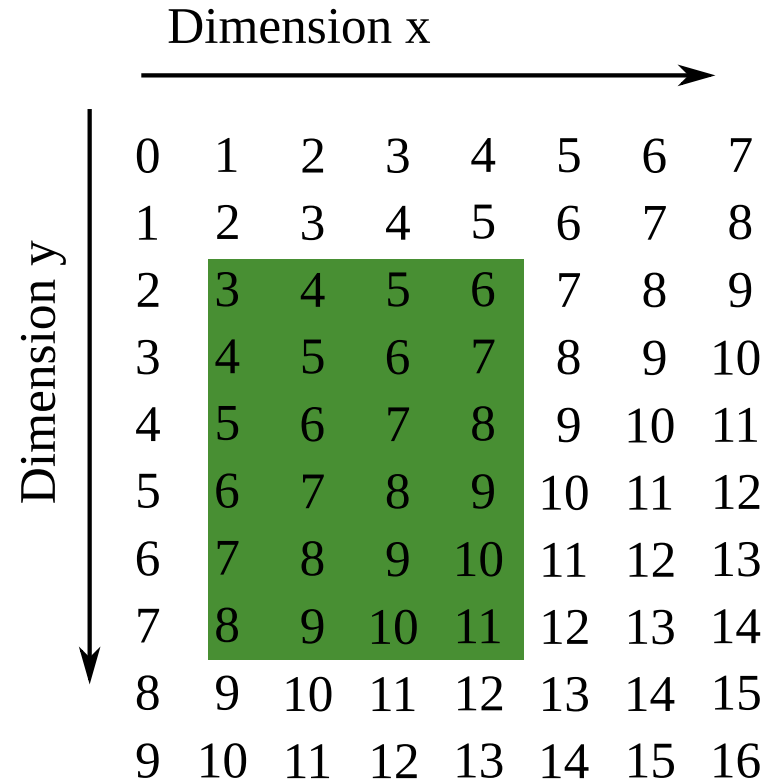
Exercise 1

Please draw the elements selected by the start, stride, count, block set below

	Dimension x →							
Dimension y ↓	0	1	2	3	4	5	6	7
	1	2	3	4	5	6	7	8
	2	3	4	5	6	7	8	9
	3	4	5	6	7	8	9	10
	4	5	6	7	8	9	10	11
	5	6	7	8	9	10	11	12
	6	7	8	9	10	11	12	13
	7	8	9	10	11	12	13	14
	8	9	10	11	12	13	14	15
	9	10	11	12	13	14	15	16

```
start[0] = 2;  start[1] = 1;  
count[0] = 6;  count[1] = 4;
```


Solution 1



Dimension x

Dimension y

0	1	2	3	4	5	6	7
1	2	3	4	5	6	7	8
2	3	4	5	6	7	8	9
3	4	5	6	7	8	9	10
4	5	6	7	8	9	10	11
5	6	7	8	9	10	11	12
6	7	8	9	10	11	12	13
7	8	9	10	11	12	13	14
8	9	10	11	12	13	14	15
9	10	11	12	13	14	15	16

```
start[0] = 2;  start[1] = 1;  
count[0] = 6;  count[1] = 4;
```

Exercise 2

Please draw the elements selected by the start, stride, count, block set below

Dimension x

Dimension y

0	1	2	3	4	5	6	7
1	2	3	4	5	6	7	8
2	3	4	5	6	7	8	9
3	4	5	6	7	8	9	10
4	5	6	7	8	9	10	11
5	6	7	8	9	10	11	12
6	7	8	9	10	11	12	13
7	8	9	10	11	12	13	14
8	9	10	11	12	13	14	15
9	10	11	12	13	14	15	16

```
start[0] = 2;  start[1] = 1;
count[0] = 1;  count[1] = 1;
block[0] = 6;  block[1] = 4;
```

Solution 2

Dimension x

Dimension y

0	1	2	3	4	5	6	7
1	2	3	4	5	6	7	8
2	3	4	5	6	7	8	9
3	4	5	6	7	8	9	10
4	5	6	7	8	9	10	11
5	6	7	8	9	10	11	12
6	7	8	9	10	11	12	13
7	8	9	10	11	12	13	14
8	9	10	11	12	13	14	15
9	10	11	12	13	14	15	16

```
start[0] = 2;  start[1] = 1;  
count[0] = 1;  count[1] = 1;  
block[0] = 6;  block[1] = 4;
```

Exercise 3

Please draw the elements selected by the start, stride, count, block set below

A 10x8 grid of numbers from 0 to 16. The horizontal axis is labeled 'Dimension x' with an arrow pointing right. The vertical axis is labeled 'Dimension y' with an arrow pointing down. The grid contains the following values:

0	1	2	3	4	5	6	7
1	2	3	4	5	6	7	8
2	3	4	5	6	7	8	9
3	4	5	6	7	8	9	10
4	5	6	7	8	9	10	11
5	6	7	8	9	10	11	12
6	7	8	9	10	11	12	13
7	8	9	10	11	12	13	14
8	9	10	11	12	13	14	15
9	10	11	12	13	14	15	16

```
start[0] = 2;  start[1] = 1;
count[0] = 3;  count[1] = 2;
stride[0] = 2; stride[1] = 2;
block[0] = 2;  block[1] = 2;
```

Solution 3

	Dimension x →							
Dimension y ↓	0	1	2	3	4	5	6	7
	1	2	3	4	5	6	7	8
	2	3	4	5	6	7	8	9
	3	4	5	6	7	8	9	10
	4	5	6	7	8	9	10	11
	5	6	7	8	9	10	11	12
	6	7	8	9	10	11	12	13
	7	8	9	10	11	12	13	14
	8	9	10	11	12	13	14	15
	9	10	11	12	13	14	15	16

```
start[0] = 2;  start[1] = 1;
count[0] = 3;  count[1] = 2;
stride[0] = 2;  stride[1] = 2;
block[0] = 2;  block[1] = 2;
```

What is a dataspace ?


Dataspace Objects

- Null dataspaces
- Scalar dataspaces
- Simple dataspaces
 - rank or number of dimensions
 - current size
 - maximum size (can be unlimited)
- Complex dataspaces to select data


What is a dataspace ?

Dataspaces come into play:

- for performing partial IO
- to describe the shape of HDF5 dataset



What is a dataspace for ?



How to play with dataspace

```
hid_t space_id;  
hsize_t dims[2], start[2], count[2];  
hsize_t *stride=NULL, *block=NULL;  
  
dims[0] = ny; dims[1] = nx;  
start[0] = 2; start[1] = 1;  
count[0] = 6; count[1] = 4;  
  
space_id = H5Screate_simple(2, dims, NULL);  
  
status = H5Sselect_hyperslab(space_id, H5S_SELECT_SET, start, \  
stride, count, block);
```

How to play with dataspace

- *space_id* is modified by *H5Sselect_hyperslab*, so it must exist
- *start*, *stride*, *count*, *block* arrays must be at least the same size as the rank of *space_id* dataspace
- H5S_SELECT_SET replaces the existing selection with the parameters from this call.
- Other operations : H5S_SELECT_OR, AND, XOR, NOTB and NOTA
- *stride*, *block* arrays are considered as 1 if NULL is passed

Using dataspace during a partial IO

```
status = H5Sselect_hyperslab(space_id_mem, H5S_SELECT_SET, \
start_mem, stride_mem, count_mem, block_mem);

status = H5Sselect_hyperslab(space_id_disk, H5S_SELECT_SET, \
start_disk, stride_disk, count_disk, block_disk);

status = H5Dwrite(dataset, H5T_NATIVE_INT, space_id_mem, \
                 space_id_disk, H5P_DEFAULT, data);
```

- The two dataspace can describe non contiguous data and can be of different dimension
- But the number of elements must match

HDF5 command line tools

- HDF5 files are non ASCII files
- non human readable files
- Tools provided to manipulate and get information contained in HDF5 files
- Three main ones: **h5ls**, **h5dump**, **h5diff**

Hands on HDF5

- [hands-on on github](#)
- Just use the one prepared for this training by

```
cp -r /home/prace/HDF5_hands-on .
```

Hands on HDF5

